

Locally Boosted Graph Aggregation and Community Detection

Submitted for blind review

Abstract—Learning the right graph representation from noisy, multi-source data has garnered significant interest in recent years. A central tenet of this problem is relational learning. Here the objective is to incorporate the partial information each data source gives us in a way that captures the true underlying relationships. To address this challenge, we present a general, boosting-inspired framework for combining weak evidence of entity associations into a robust similarity metric. We explore the extent to which different local quality measurements yield graph representations that are suitable for community detection. We present empirical results on a variety of datasets demonstrating the utility of this framework, especially with respect to real datasets where noise and scale present serious challenges. Finally, we prove a convergence theorem and outline future research into other application domains.

I. INTRODUCTION

When studying networks, the data used to define nodes and edges often come from multiple sources. These sources can be noisy and ambiguously useful, and the process of combining them into a single graph representation is critically important. For example, when studying communities in social networks the data that indicate membership in the same community are plentiful: communication, physical proximity, reported friendship, etc. Each data source carries different information about the underlying social structure, and each may accurately represent only some of the individuals. Some groups of friends communicate through Facebook and others via Instagram, etc. The best way to aggregate this information is unclear, and the choice of graph representation significantly impacts the performance of subsequent data mining algorithms [17], [16], [30], [9], [27]. Further complicating matters, the quality of the aggregated graph depends on the application domain. Community detection is easiest with a graph representation that discards cross-community edges, but to predict the spread of a virus one must keep crucial cross-community conduits. Suitable graph representations for these tasks may use the same data but are different enough to need different aggregation techniques.

Even though the impact of the graph representation on subsequent analysis has been widely studied, there are few techniques for learning good graph representations. Aggregation is often ad-hoc in practice, making it difficult to compare algorithms within the same domain using different data sources. The need for rigorous approaches to graph representation learning is even more apparent with big data.

In this paper, we present a graph aggregation framework designed to make the process of learning a useful underlying graph representation rigorous with respect to application specific requirements. Our framework is called *Locally Boosted*

Graph Aggregation (LBGA). LBGA extracts the application-specific aspects of the learning objective as an event A representing an operation on the graph (e.g. a clustering algorithm, a random walk, etc.) and a local quality measure q . Using these, the framework simulates a reward system that promotes the presence of good edges and the absence of bad edges, in a fashion inspired by boosting.

We demonstrate LBGA applied to community detection. In this context the goal of graph representation learning is to aggregate the different data sources into a single graph which makes the true community structure easy to detect. LBGA evaluates the graph data locally, so that it can choose the data sources which most accurately represent the local structure of communities observed in real networks [2], [23]. In the absence of ground truth knowledge or one efficiently computable measure that can capture true community quality, LBGA relies on the pair of a graph clustering algorithm A and a local clustering metric q as an evaluation proxy. We show through empirical analysis that our algorithm can learn a high-quality global representation guided by the local quality measures considered.

We make the following contributions:

- 1) We present a graph aggregation framework that learns a useful graph representation with respect to an application requiring only a local heuristic measure of quality.
- 2) Our framework is stochastic and incorporates both edge and non-edge information, making it robust and suitable for sparse and noisy networks.
- 3) We demonstrate the success of an algorithm implementing the framework for community detection, testing it against both synthetic data and real-world data sets.
- 4) We perform sensitivity and scalability analyses of our algorithm, showing that the algorithm scales linearly in the number of edges and is robust enough to handle large, noisy graphs.
- 5) We prove a convergence theorem for our framework and suggest the next steps in proving performance guarantees.

We emphasize that while this paper specifically addresses community detection, the central focus here is on aggregating multiple noisy, potentially adversarial graphs into a single graph. Even though the definition of what makes an edge qualitative depends on the application, the core piece of our algorithm, the edge reward mechanism and the graph aggregation step, are application agnostic.

This paper is organized as follows. In Section II we review related literature. In Section III we discuss in detail the LBGA framework. In Section IV we present the experimental analysis and results. In Section V we discuss sensitivity to noise and

scalability. In Section VI we prove our convergence theorem, and in Section VII we discuss future work.

II. RELATED WORK

A. Representation learning and clustering

Representation learning has garnered a lot of interest and research in recent years. Its goal is to introduce more rigor to the often ad-hoc practices of transforming raw, noisy, multi-source data into inputs for data mining and machine learning algorithms. Within this area, representation learning of graph-based data includes modeling decisions about the nodes of the graph, the edges, as well as the critical features that characterize them both.

In this context, Rossi et al. [35] discuss transformations to heterogeneous graphs (graphs with multiple node types and/or multiple edge types) in order to improve the quality of a learning algorithm. Within their taxonomy, our work falls under the link interpretation and link re-weighting algorithms [40], [18]. Our setting is different because we explicitly allow different edge types between the same pair of vertices. Also, our approach is stochastic, which we find necessary for learning a robust representation and weeding out noise.

Clustering in multi-edge graphs [32], [38], [37], [28], [6], [22] is another area with close connections to our work. A common thread among these existing approaches is clustering by leveraging shared information across different graph representations of the same data. These approaches do not address scenarios where the information provided by the different sources is complementary or the overlap is scarce. In contrast, our approach iteratively selects those edge sources that lead to better clustering quality, independently of disagreement across the different features. [34], [10] present approaches for identifying the right graph aggregation, given a complete ground truth clustering, or a portion of it (i.e.: the cluster assignment is known only for a subset of the vertices in the graph). Our framework requires no such knowledge, but we do use ground truth to validate our experiments on synthetic data (Section IV-C). Balcan and Blum define in [4], [5] a list of intuitive properties a similarity function needs to have in order to be able to cluster well. However, testing whether a similarity function has the discussed properties is NP-hard, and often dependent on having ground truth available. Our model instead uses an efficiently computable heuristic as a rough guide.

B. Boosting and bandits

Our framework departs from previous work most visibly through its algorithmic inspirations, namely boosting [36] and bandit learning [7]. Neither framework applies directly to our problem, but they suggest a natural algorithm whose simplifications have measurable convergence properties.

In boosting, one assumes the existence of a *weak classifier* whose performance is slightly better than random. In a landmark paper [36], Schapire showed how to combine weak classifiers into a PAC-learner by a majority voting scheme. One can consider different graph data sources as weak learners, and

ask whether one can “boost” them to a good graph. Unfortunately, our problem setting does not allow pure boosting for two reasons: input graphs can be pure noise or adversarially bad and hence are not weak learners; and boosting requires access to ground truth labels. Even with reliable input, the application domain may have no accepted measure of quality.

Ideas from bandit learning compensate for these problems. In bandit learning an algorithm receives rewards as it explores a set of actions, and the goal is to minimize some notion of regret in hindsight. The basic model has many variants, but two central extensions in the literature are expert advice and adversaries. Experts are functions suggesting what action to take in each round (which are arbitrary). The adversarial setting involves an omniscient adversary who sets the experts and rewards so as to maximize regret. In particular, rewards can vary across rounds.

Using these ideas, we set up a reward system based on the given application and use stochastic weight update techniques to learn a graph representation. In our setting we only care if the aggregate graph is good at the end, while bandit learning often seeks to maximize cumulative rewards during learning. There are bandit settings that only care about the final result (e.g., pure exploration [8]), but to the best of our knowledge they do not apply to our problem.

The primary technique we adapt from bandits and boosting is the Multiplicative Weights Update Algorithm (MWUA). See [3] for an overview and an extensive list of successful applications. The algorithm maintains a weight for each element x_j of a finite set X . In rounds, an element x_i is chosen by sampling proportionally to the weights, a reward $q_{t,i}$ is received, and the weight for x_i is multiplied or divided by $(1 + \varepsilon q_{t,i})$, for some parameter $\varepsilon > 0$. After many rounds, the elements with the highest weight are deemed the best and used for whatever purpose needed. Next, we describe how this algorithm is adapted to graph aggregation.

III. THE LOCALLY BOOSTED GRAPH AGGREGATION FRAMEWORK

The Locally Boosted Graph Aggregation framework (LBGA) can succinctly be described as running MWUA for each possible edge, forming a candidate graph representation G_t in each round by sampling from all edge distributions, and computing local rewards on G_t to update the weights for the next round. Over time G_t stabilizes and we produce it as output. The remainder of this section expands the details of this sketch and our specific algorithm implementing it.

A. Framework details

Let H_1, \dots, H_m be a set of unweighted, undirected graphs defined on the same vertex set V . We think of each H_i as “expert advice” suggesting for a pair of vertices $u, v \in V$ whether to include edge (u, v) or not. Our goal is to combine the information present in the H_i to produce a global graph representation G^* suitable for a given application.

We present LBGA in the context of community detection, noting generalizations. Each round has four parts: producing the aggregate candidate graph G_t , computing a clustering A

for use in measuring the quality of G_t , computing the local quality of each edge, and using the quality values to update the weights for the edges. After some number of rounds T , the process ends and we produce $G^* = G_T$.

Aggregated Candidate Graph G_t : In each round produce a graph G_t as follows. Maintain a weight $w_{u,v,i}$ for each graph H_i and each edge (u, v) in $H_1 \cup \dots \cup H_m$. Normalize the set of all weights for an edge $w_{u,v}$ to a probability distribution over the H_i ; thus one can sample an H_i proportionally to its weight. For each edge, sample in this way and include the edge in G_t if it is present in the drawn H_i .

Event $A(G_t)$: After the graph G_t is produced, run a clustering algorithm A on it to produce a clustering $A(G_t)$. In this paper we fix A to be the Walktrap algorithm [33], though we have observed the effectiveness of other clustering algorithms as well. In general A can be any event, and in this case we tie it to the application by making it a simple clustering algorithm.

Local quality measure: Define a *local quality measure* $q(G, e, c)$ to be a $[0, 1]$ -valued function of a graph G , an edge e of G , and a clustering c of the vertices of G . The quality of (u, v) in G_t is the “reward” for that edge, and it is used to update the weights of each input graph H_i . More precisely, the reward for (u, v) in round t is $q(G_t, (u, v), A(G_t))$.

Update Rule: Update the weights using MWUA as follows. Define two learning rate parameters $\varepsilon > 0, \nu > 0$, with the former being used to update edges from G_t that are present in H_i and the latter for edges not in H_i . In particular, suppose $q_{u,v}$ is the quality of the edge (u, v) in G_t . Then, the update rule is defined as follows:

$$w_{u,v,i} = \begin{cases} w_{u,v,i}(1 + \varepsilon q_{u,v}), & \text{if } (u, v) \in H_i \\ w_{u,v,i}(1 - \nu q_{u,v}), & \text{if } (u, v) \notin H_i. \end{cases}$$

B. Quality measures for community detection

We presently describe the two local quality measures we use for community detection. The first, which we call *Edge Consistency (EC)* asserts that edges with endpoints in the same cluster are superior to edges across clusters:

$$EC_{u,v} = \begin{cases} 1, & \text{if } c(u) = c(v) \\ -1, & \text{if } c(u) \neq c(v). \end{cases}$$

EC offers a quality metric that is inextricably tied to the performance of the chosen clustering algorithm, and as such we have found it performs poorly in practice. However, edge consistency can be combined with any quality function q to produce a “consistent” version of q . Simply evaluate q when the edge is within a cluster, and $-q$ when the edge is across clusters. Note that q can represent an algorithmic-agnostic measure of clustering quality.

As an example, we define *Neighborhood Overlap (NO)*, which uses the idea that vertices that share many neighbors are likely to be in the same community. NO declares the quality of (u, v) to be the (normalized) cardinality of the intersection of the neighborhoods of u and v , namely $NO_{u,v} = \frac{|N(u) \cap N(v)|}{|N(u) \cap N(v)| + \log(|V|)}$, where $N(x)$ represents the neighborhood of vertex x . We have also run experiments using

more conventional normalizing mechanisms, such as the Dice and Jaccard indices [12], [20]), but our neighborhood overlap metric outperforms them by at least 10% in our experiments. We argue this is due to the use of a global normalization factor, as opposed to a local one, which is what Dice and Jaccard indices use. For brevity and simplicity, we omit our results for Jaccard and Dice indices and focus on Neighborhood Overlap. In our experimental analysis (Section IV-D) we use the consistent version of *NO*, which we denote *consistentNO*.

While we demonstrate the utility of the LBGA framework by using *consistentNO*, the design of the framework is modular, in that the mechanism for rewarding the “right” edges is independent from the definition of reward. This allows us to plug in other quality metrics to guide the graph representation learning process for other applications.

Lastly, the quality of an edge is dictated by the choice of the clustering algorithm (event). In our case, walktrap merges clusters with the objective of maximizing modularity and this affects the presence of small clusters. To alleviate this, one can use a modified walktrap that outputs similarity values, or re-run LBGA on the graphs induced by the vertices belonging to the resulting clusters. The latter gives a hierarchical clustering, and LBGA can identify hierarchical community structure in this way, though we omit the details for brevity.

C. LBGA implementation

Processing every edge in every round of the LBGA framework is inefficient. Our implementation of LBGA, given by Algorithm 1, improves efficiency by fixing edges whose weights have grown so extreme so as to be picked with overwhelming or negligible probability (with probability $> 1 - \delta$ or $< \delta$ for a new parameter δ). In practice this produces a dramatic speedup on the total runtime of the algorithm, because as the algorithm learns the sampling procedure becomes substantially sublinear in the number of edges. The worst-case time complexity is the same, but we discuss practical methods to speed up LBGA in Section V-C.

In addition, our decision to penalize non-edges ($\nu > 0$) also improves runtime from the alternative ($\nu = 0$). In our experiments non-edge feedback causes G_t to converge in roughly half as many rounds as when only presence of edge is considered as indication of relational structure. Moreover, the algorithm is stable to minor variations in ε and δ . Indeed, the theoretical guarantees of MWUA in general and Section VI suggest that for any fixed $0 < \varepsilon, \delta < 1/2$ the algorithm converges in at most $O(\log(n/\delta))$ rounds.

We also note that Algorithm 1 stays inside the “boundaries” determined by the input graphs H_i . It never considers edges that are not suggested by *some* H_i , nor does it reject an edge suggested by all H_i . Thus, when we discuss sparsity of our algorithm’s output in our experiments, we mean with respect to the number of edges in the union of the input graphs.

IV. EXPERIMENTAL ANALYSIS

We describe the datasets used for analysis and provide quantitative results for the performance of Algorithm 1.

Data: Unweighted graphs H_1, \dots, H_m on the same vertex set V , a clustering algorithm A , a local quality metric q , three parameters $0 < \varepsilon, \nu, \delta < 1/2$

Result: A graph G

Initialize a vector $\mathbf{w}_{u,v} = \mathbf{1}$ for all $u \neq v \in V$

Let U be the edge set of $H_1 \cup \dots \cup H_m$

Let $G_{\text{learned}} = (V, \emptyset)$

while $|U| > 0$ **do**

 Let G be a copy of G_{learned}

for $(u, v) \in U$ **do**

 Let $p_{u,v} = \frac{\sum_i w_{u,v,i} 1_{\{(u,v) \in H_i\}}}{\sum_i w_{u,v,i}}$

 Flip a coin with bias $p_{u,v}$

 If heads, include (u, v) in G .

end

 Cluster G using A

for $(u, v) \in U$ **do**

 Set $p = q(G, A(G), (u, v))$

for $i = 1, \dots, m$ **do**

if $(u, v) \in H_i$ **then**

 Set $w_{u,v,i} = w_{u,v,i}(1 + \varepsilon p)$

else

 Set $w_{u,v,i} = w_{u,v,i}(1 - \nu p)$

end

end

 Let $p_{u,v} = \frac{\sum_i w_{u,v,i} 1_{\{(u,v) \in H_i\}}}{\sum_i w_{u,v,i}}$

if $p_{u,v} > 1 - \delta$ **then**

 Add (u, v) to G_{learned} , remove it from U

end

if $p_{u,v} < \delta$ **then**

 Remove (u, v) from U

end

end

end

Output G

Algorithm 1: Optimized implementation of LBGA. Note that 1_E denotes the characteristic function of the event E .

A. Synthetic datasets

Our primary synthetic data model is the stochastic block model [39], commonly used to model explicit community structure. We construct a probability distribution $G(\mathbf{n}, B)$ over graphs as follows. Given a number n of vertices and a list of cluster (block) sizes $\mathbf{n} = \{n_1, \dots, n_k\}$ such that $n = \sum_i n_i$, we partition the n vertices into k blocks $\{b_1, \dots, b_k\}$, $|b_i| = n_i$. We declare that the probability of an edge occurring between a vertex in block b_i and block b_j is given by the (i, j) entry of a k -by- k matrix B .

A reformulation of the stochastic block model in the multimodal setting considers how strongly each graph source represents each community in data. In the most local case, each graph source represents well only one community. We can model this as follows. Fix some probability p_i of an edge occurring within a community and r_i for an edge occurring across communities ($p_i \gg r_i$). Then draw the input graph H_i from the distribution $G(\mathbf{n}, B_i)$, where B_i has p_i in its (i, i) entry and r_i everywhere else. For example, if there are $m = 2$ communities the probability matrices would be

$$B_1 = \begin{pmatrix} p_1 & r_1 \\ r_1 & r_1 \end{pmatrix}, \quad B_2 = \begin{pmatrix} r_2 & r_2 \\ r_2 & p_2 \end{pmatrix}.$$

We call this model the *local stochastic block model* (LSBM).

Dataset	Parameters
GSBM-4	$m = k = 4, n_i = 125, p_1 = 0.1625,$ $p_2 = 0.125, p_3 = 0.125, p_4 = 0.0875, r_i = 0.05$
GSBM-5	$m = k = 4, n_i = 125, p_1 = 0.15, p_2 = 0.1,$ $p_3 = p_4 = 0.05, r_i = 0.05, i = 1, \dots, m$
LSBM-1	$m = k = 4, n_i = 125, p_i = 0.2, r_i = 0.05$
LSBM-2	$m = k = 4, n_i = 125, p_i = 0.3, r_i = 0.05$
LSBM-3	$m = 5, k = 4, n_i = 125, p_i = 0.3, r_i = 0.05,$ $i = 1, \dots, m, p_5 = r_5 = 0.01$
ER only	$m = 4, p_i = r_i = 0.01$
DBLP	$n = 3153, m = 2$
RealityMining	$n = 90, m = 6$
Enron	$n = 145, m = 2, \alpha = 0.9$

TABLE I: Description of datasets analyzed. Total number of vertices in each synthetic source graph is $n = 500$. m is the number of graph sources. k is the number of clusters. n_i represents number of vertices in cluster i . p_i and r_i represent the within- and across-cluster edge probability for each the m graph sources.

We can further vary the number of communities each graph source strongly represents, with the other extreme being the case where each graph source has a global contribution toward the overall community structure. We refer to this case as the *global stochastic block model* (GSBM).

Finally, we consider the case of the *Erdős-Rényi random graph* [14], where any two vertices have equal probability of being connected. This model provides an example of a graph with no community structure. Note that the ER model is a special case of LSMB with $p = r$. We also consider cases where an ER model is injected into block model instances in order to capture a range of structure and noise combinations.

B. Real datasets

1) *DBLP*: DBLP [24] is a comprehensive online database documenting research in computer science. We extracted the subset of the DBLP database corresponding to researchers who have published at two conferences: the Symposium on the Theory of Computing (STOC), and the Symposium on Foundations of Computer Science (FOCS). The breadth of topics presented at these conferences implies a natural community structure organized by sub-field. Each node in the DBLP graph represents an author, and we use two graphs on this vertex set: the *co-authorship* graph and the *title similarity* graph. For the latter, we add an edge between two author vertices if any of their paper titles contain at least three words in common (excluding stop words), and the weight of this edge is the number of such pairs of papers. We considered 5234 papers across 3153 researchers.

2) *RealityMining*: Our second dataset is RealityMining [13], a 9-month experiment in 2004 which tracked a group of 90 individuals at MIT via sensors in their cell phones. The individuals were either associated with the MIT Media Lab or the Sloan Business School, and there is a natural corresponding community structure. The data collected include voice calls, bluetooth scan events at five-minute intervals, cell tower usage, and self-reported friendship and proximity data.

We used the subset of subjects participating between 2004-09-23 19:00:00 and 2005-01-07 18:00:00 (UTC-05:00), for a total of 3354 call events, 786301 cell tower transition events,

and 689025 bluetooth scan events. The nodes in our graphs represent individuals in the study. Weighted edges correspond to the total duration of voice calls, the total amount of time two individuals used the same cell tower, the total number of bluetooth events, and the results of the friendship/proximity surveys for a total of six graphs.

3) *Enron*: Our final dataset is the Enron email dataset [21], [1], a well-studied corpus of over 600,000 emails sent between 145 employees of the Enron Corporation in the early 2000's. We produced two graphs from the Enron data, one for peer-to-peer email communication and one for topic similarity in the email content.

In both graphs the vertices are individuals. In the email graph, the edges are weighted by the number of emails sent between the individuals in question. We used the Mallet package [26] to generate the LDA topic model for the content of Enron email data. We aggregated into one document all the email content sent by each of the Enron employees considered in the email link graph. Each document and therefore each sender is represented by 60 topics. We measure cosine distance of the topic vectors of individuals, and considered an edge as present if the cosine distance was above a specified threshold value α .¹

Table I contains a summary of all the datasets used for the experimental analysis and their parameters.

C. Validation procedure

In our work, the optimality of the graph representation is closely coupled with the quality of community structure captured by the representation. This gives us several ways of evaluating the quality of the results produced by our algorithm. We consider notions of quality reflected at different levels: the quality of cluster assignment, the quality of graph representation, and the quality of graph source weighting.

Quality of Cluster Assignment: Since the output of LBGA is a graph, we use the walktrap clustering algorithm to extract communities for analysis. We then compare these communities to the ground truth clustering, when it is available, or else to the known features of the datasets. We use the Normalized Mutual Information (NMI) measure [11] to capture how well the ground truth clustering overlaps with the clustering on the graph representation output from our algorithm.

Quality of Graph Representation: An ideal graph representation that contains community structure would consist of disjoint cliques or near-cliques corresponding to the communities. As we illustrate in Section IV-D, an optimal graph representation can do better than just produce a perfect clustering. It can also remove cross-community edges and produce a sparser representation, which is what our algorithm does. We use two measures of clusteriness to capture this notion of graph representation quality. Modularity [31] is a popular measure that compares a given graph and clustering to a null model. Conductance [23], [19] measures relative sparsity of cuts in the graph. Note that *higher* modularity scores and *lower* conductance scores signify stronger community structure. Both

modularity and conductance are well-known and often offer complimentary information about the quality of communities.

We note two extreme graph representation cases, the empty graph which is perfectly modular in a degenerate sense, and the union graph which is a trivial aggregation. To signal these cases in our results, we display the *sparsity* of the produced graph G^* , defined as the fraction of edges in G^* out of the total set of edges in all input graphs.

Quality of Graph Source Weighting: the quality of the aggregation process is captured by the right weighting of individual edge sources. Edge sources (input graphs) that are more influential in uncovering the underlying community structure have higher weights on average. Similarly, edge types that contribute equally should have equal weights, and edge types with no underlying structure should have low weights.

D. Experimental results

Table II contains the numerical results of our experiments. As a baseline, we computed the modularity and conductance values of the union of the input graphs with respect to the ground truth (for synthetic) or the Walktrap clusterings (for the real world). For synthetic examples we also compare with GraphFuse [32].

Overall, we find that LBGA converges to graphs of both high modularity and low conductance score. It also generates graph representations that induce correct clusterings in almost all cases where some sort of ground truth is known, the challenging case being when SNR is low.

1) *Synthetic*: For illustration, we show in Figure 1 the performance of Algorithm 1 when *consistentNO* is used as a local quality metric and LSBM-3 (see Table I for details) is used to generate the input graphs. Note that the algorithm converges quickly to a graph which results in a perfect clustering as measured by NMI. We also plot the modularity of the resulting graph produced in each round, seeing that it far exceeds the “baseline” modularity of the union of the input graphs. This tells us the learning algorithm is able to discard the noisy edges in the model. Finally, we plot the number of edges in the graph produced in each round, and the average vertex-pair weight for each input graph. This verifies that our algorithm complies with our edge-type weighting and sparsity requirements. Indeed, the algorithm produces a relatively sparse graph, using about 40% of the total edges available and weights edges from the Erdős-Rényi source appropriately. Our algorithm hence achieves a superior graph than the union, while preserving the underlying community structure so as to be amenable to clustering.

The results for the other synthetic datasets are similar and summarized in Table II. We note that our algorithm's performance degrades when the noise becomes too high. In Section V-B we analyze the signal to noise ratio in the synthetic data sets more closely.

2) *DBLP*: The data of Table II shows our algorithm converges to a DBLP graph of modularity exceeding that of the union graph and using significantly fewer edges. Our algorithm selects title similarity as having more influence in recovering communities for the STOC/FOCS conferences. Researchers

¹We experimented with various threshold values, and we discuss this in Section IV-D3.




Fig. 1: Graph representation learning for LSMB-3. The LPGA parameters are $\varepsilon = \nu = 0.2$, $\delta = 0.05$. Plots in order top to bottom: 1. NMI of $A(G_t)$ with the ground truth clustering, 2. modularity of G_t w.r.t $A(G_t)$, with the horizontal line showing the modularity of the union of the input graphs w.r.t. ground truth, 3. the number of edges in G_t , 4. the average probability weight (quality) of vertex pairs for H_i . The Erdős-Rényi graph converges to low weight by round 300, even though it is initially favored. This is evidence of LPGA’s ability to recover from initial bad luck.




Fig. 2: LPGA with consistentNO on a dataset of Erdős-Rényi random graphs.

attending these conferences represent a small community as a whole, with many of them sharing co-authorship on papers with diverse topics. In this sense, it is not surprising that title similarity serves as a better proxy for capturing the more pronounced division along topics. We have also manually inspected the resulting clusters, and most appear organized both by membership and coauthorship. Take, for example,




Fig. 3: Statistics about the aggregate graph produced by LPGA after 500 rounds on a suite of 4 Erdős-Rényi random graphs on 500 nodes and varying edge probability p .

Mikko Koivisto, Thore Husfeldt, Petteri Kaski, and Andreas Björklund. They have together coauthored over 15 papers in combinatorial optimization, and naturally fall within the same small coauthorship cluster. However, the title similarity graph alone yields around 1500 clusters, and these researchers are split across two clusters because of the differences in their non-coauthored work. They fall in the same cluster in the LPGA-aggregated graph, and the cluster is larger, including well-known researchers who are either coauthors with some of the four or have done much work in the same field.

3) RealityMining & Enron: The final aggregated RealityMining graph constructed by LPGA contains two dense clusters corresponding exactly to the MIT Media Lab and the Sloan Business School, with only three edges crossing the cut. In addition, this graph uses only 63.5% of the total edges available (see Table II).

For Enron, the data of Table II shows that *consistentNO* achieves a graph representation with higher modularity, lower conductance and better sparsity when compared to the baseline. Figure 4 shows a clear community structure, and there the smaller clusters correspond to lower-level employees while the higher level managers reside in the bigger clusters. Additionally, there was a known fantasy sports community within the network, and all of these individuals fall within a single cluster in the graph output by LPGA [25].

We also investigated the effect of changing the threshold value α for considering an edge in the Enron topic graph. A value of less than $\alpha = 0.7$ always produced two large dense clusters with many noisy edges between them (akin to the two large clusters in Figure 4). In our experiments we used $\alpha = 0.9$, although values of α as high as 0.95 gave qualitatively similar results.

We notice that the modularity values for RealityMining and the Enron data set are significantly smaller after passing through LPGA when compared to the baseline union values.




Fig. 4: Left: the results of LBGA on the Enron dataset. Right: the input graph of topic models thresholded at 0.9. LBGA was run with *consistentNO*, $\nu = \varepsilon = 0.2$, $\delta = 0.05$

We argue that this is due to the relatively small clusters produced by LBGA, as it is a known shortcoming that modularity is not an accurate measure when the communities are small [15]. Indeed, when conductance is used as a clustering quality measure LBGA significantly outperforms the baseline union aggregation.

E. Comparison with GraphFuse

We compare LBGA with GraphFuse [32], a multi-graph clustering algorithm that falls under the category of tensor-based clustering [22], [38]. GraphFuse computes the clustering based on the CP decomposition of the tensor formed by appending the adjacency matrices of the different graph sources. It differs fundamentally from LBGA in that it does not produce an aggregated graph representation. We use NMI with the ground truth as the performance measure. For the comparison analysis we have only considered the synthetic datasets where the notion of ground truth is clear and avoided the real datasets where the notion of ground truth is subjective. Table II contains the comparison results.

We find that *consistentNO* outperforms GraphFuse on both the global block models and the lower-noise local block models LSM-2 and LSM-3. LBGA also produces very sparse representations that may be useful for future analysis, while GraphFuse only produces a clustering.

V. ROBUSTNESS AND SCALING

A. LBGA does not boost noise

We consider whether LBGA falsely boosts noise to “find” community structure where none is present. While this depends crucially on the choice of quality function and event, we find that for *consistentNO* and walktrap clustering LBGA does not falsely boost noise. Indeed, we even claim that LBGA recovers from initially poor luck, which agrees with the theoretical adversarial bandit foundations. As evidence we consider Figures 1 and 2, which respectively depict the behavior of LBGA on a local stochastic block model with an extra Erdős-Rényi graph and a dataset comprised entirely of Erdős-Rényi random graphs. The assumption is that Erdős-Rényi graphs have no community structure.

In Figure 1 we see that initially LBGA weights the Erdős-Rényi graph (the pink line) higher than the graphs from the LSM model, but by round 300 it has recovered and ultimately

produces the ground truth clustering. In Figure 2 we see that LBGA produces aggregate graphs whose maximal modularity values and minimum conductance values are poor, and the final representation is very sparse. In both of these cases we see that LBGA largely discards noise.

Finally, we observe the phase transition for a dataset of Erdős-Rényi random graphs (on $n = 500$ vertices) as the probability of an edge increases in Figure 3. We see a clear phase transition around $p = 0.3$, before which modularity values are low and conductance values are high. This provides evidence that LBGA with *consistentNO* does not falsely boost sparse amounts of noise, but will err on datasets with very high levels of noise.

B. Sensitivity analysis

We analyze the sensitivity of LBGA to noise. In Figure 5 we display performance as measured by NMI when the graph inputs are LSM models across different intra-cluster probability values p_i and varying SNR values. We make the following general observations. The algorithm is consistent in that as the noise rate r_i increases, the NMI values do deteriorate as expected. The algorithm both reaches higher quality and maintains the quality longer for denser graphs, which again is consistent with our expectations. At a signal to noise ratio of 2 or less, the NMI is bad for *consistentNO* and all choices of p_i . The sharp drop in quality is related to the well-known phase transition in the community detection problem [29]. Therefore, the distance from the theoretical detectability bound on community detection could be used as an agnostic measure to gauge the usefulness of LBGA with a particular quality metric.




Fig. 5: Performance of LBGA (measured by NMI) as a function of SNR for the LSM model with different probabilities p_i for *consistentNO*.

C. Scalability analysis

We run a larger version of the LSM-2 model ($m = k = 10, p_i = 0.3, r_i = 0.05$) to illustrate how LBGA scales for larger graphs. As shown in Figure 6, LBGA scales linearly with the number of edges. Given that real-world graphs are usually sparse, this scaling behavior makes LBGA computationally suitable for large graphs.

Union Graph				GraphFuse	LBGA: ConsistentNO					
Dataset	Modularity	Conductance	NMI	NMI	Modularity	Conductance	NMI	Sparsity	Edge Type	Weights
GSBM-4	0.178	10.678	1.000	0.716	0.739	0.084	1.000	0.433	(3.46,2.92,3.08,2.43)	
GSBM-5	0.093	15.368	0.636	0.616	0.727	2.121	0.619	0.235	(2.02,1.75,1.33,1.32)	
LSBM-1	0.103	14.725	0.724	0.686	0.568	3.322	0.536	0.274	(1.86,1.91,1.87,1.92)	
LSBM-2	0.166	11.233	0.992	0.760	0.740	0.084	0.992	0.420	(3.01,2.89,3.04,2.98)	
LSBM-3	0.166	11.216	1.000	0.779	0.737	0.104	1.000	0.422	(2.99,3.01,2.95,2.97)	
ER only	-0.002	24.729	-	-	0.193	112.947	-	0.230	(0.251,0.253,0.248,0.247)	
DBLP	0.386	1368.859	-	-	0.695	159.286	-	0.632	(0.432,0.568)	
RealityMining	0.452	70.314	-	-	0.246	0	-	0.646	(0.365,0.091,0.198,0.115,0.115,0.115)	
Enron	0.559	134.572	-	-	0.444	0.594	-	0.631	(0.390,0.610)	

TABLE II: LBGA performance results, compared to GraphFuse and a baseline union aggregation. All datasets in this table were run with *consistentNO* using $\varepsilon = \nu = 0.2$, $\delta = 0.05$. Union modularity and conductance for real datasets was computed with the walktrap clustering. The order of edge type weights for the real datasets are: DBLP (coauthorship, title similarity); RealityMining (bluetooth, phone calls, cell tower proximity, reported friendship, in-lab proximity, out-lab proximity); Enron (email, topic similarity).

LBGA was designed and implemented with scaling in mind, and the process of fixing vertex-pairs as they are learned is largely the reason for the nice scaling properties of Algorithm 1. Moreover, the parameter δ encapsulates a trade-off between runtime and accuracy. Should linear scaling be insufficient, LBGA’s design allows for additional modifications to improve scalability. For example: quality functions that are sufficiently local allow one to parallelize the weight update step; one could sample a sublinear number of edges in each round and only update weights within the subgraph; one might specify a set of “seeds” (vertices of interest), and sample edges local to those vertices. While these methods are currently just potential directions for future work, there is clearly much potential for further improvements of LBGA’s computational efficiency.




Fig. 6: The runtime of LBGA.

VI. A CONVERGENCE THEOREM

In this section we provide a theorem on the convergence rate of LBGA. The result can be interpreted as follows: if the quality function is sufficiently knowledgeable about the application domain, then LBGA will converge to the “right” graph representation. Specifically, we assume the quality function q is a probabilistic oracle for the ground-truth clustering, which can be interpreted as a local weak learner.

We first set up some notation. Let H_1, \dots, H_N be graphs on the same vertex set V . Let G^* be a target graph. For clustering we assume G^* is a disjoint union of the clusters, but the theorem applies to any target graph. Call a graph H_i *good* for a vertex pair (u, v) if (u, v) is an edge of both H_i and G^* , or (u, v) is an edge of neither H_i nor G^* . Otherwise call H_i *bad* for (u, v) . For $0 \leq \eta < 1/2$ define the *noisy oracle* Q_η so that Q_η is correct with probability $1/2 + \eta$ and incorrect with probability $1/2 - \eta$. In formulas,

$$\begin{aligned}
\Pr[Q_\eta(u, v) = 1 \mid (u, v) \in E(G^*)] &= \frac{1}{2} + \eta, \\
\Pr[Q_\eta(u, v) = -1 \mid (u, v) \in E(G^*)] &= \frac{1}{2} - \eta, \\
\Pr[Q_\eta(u, v) = 1 \mid (u, v) \notin E(G^*)] &= \frac{1}{2} - \eta, \\
\Pr[Q_\eta(u, v) = -1 \mid (u, v) \notin E(G^*)] &= \frac{1}{2} + \eta.
\end{aligned}$$

The multiplicative update for good graphs is the opposite of that for bad graphs (regardless of the value of Q_η). Thus, we can remove the conditioning, and define q_η to be a $\{-1, 1\}$ -valued random variable which is 1 with probability $1/2 + \eta$. Without loss of generality, the update rule in LBGA can be expressed by the multiplicative factor of $(1 + q_\eta \varepsilon)$ for good graphs and $(1 - q_\eta \varepsilon)$ for bad graphs. Further note that $\mathbb{E}[q_\eta] = 2\eta$ for good graphs and -2η for bad graphs.

Fix a vertex pair (u, v) . Let $w_{u,v,i}$ be the weight of (u, v) with respect to H_i . Define $w_{u,v,\text{good}} = \sum_{\text{good } H_i} w_i$

to be the sum of the weights of the good graphs for (u, v) , and $w_{u,v,\text{bad}} = \sum_{\text{bad } H_i} w_i$ the sum of the bad weights. Call $p_{u,v,\text{good}}, p_{u,v,\text{bad}}$ the corresponding probability weights of picking a good or bad input graph. When u, v are clear from the context, we will omit them from the subscripts and simply write w_{good} , etc. We add a superscript t to any quantity that changes over rounds to denote which round of LBGA the quantity refers to. Our goal is to bound p_{bad}^t for $t = O(\log(1/\delta))$. Denote by $m_{\text{bad}} = w_{\text{bad}}^1, m_{\text{good}} = w_{\text{good}}^1$, the number of bad and good input graphs, respectively, with $m = m_{\text{bad}} + m_{\text{good}}$.

Theorem 1: Set $\nu = \varepsilon < 1/2$ and $\eta < 1/2$ such that $0 < \varepsilon < 4\eta$. Fix a vertex pair (u, v) and assume $m_{u,v,\text{good}} \geq 1$. Then for any $0 < \delta < 1$ and $0 < s < 1$, after $O(\log(1/\delta) + \log(1/s) + \log(m))$ rounds, the value $p_{u,v,\text{bad}}^t < \delta$ with probability at least $1 - s$.

Proof: Call q_i a sequence of random variables that are independent and identically distributed according to q_η . It follows from our observations in defining q_η and simple induction that

$$w_{\text{good}}^t = m_{\text{good}} \prod_{i=1}^{t-1} (1 + q_i \varepsilon),$$

$$w_{\text{bad}}^t = m_{\text{bad}} \prod_{i=1}^{t-1} (1 - q_i \varepsilon).$$

By the linearity of expectation we have $\mathbb{E}[w_{\text{good}}^t] = m_{\text{good}}(1 + 2\eta\varepsilon)^{t-1}$, $\mathbb{E}[w_{\text{bad}}^t] = m_{\text{bad}}(1 - 2\eta\varepsilon)^{t-1}$, and noting $q_\eta^2 = 1$, a standard computation provides the variance:

$$\text{Var}(w_{\text{bad}}^t) = m_{\text{bad}}^2 [(1 - 4\eta\varepsilon + \varepsilon^2)^{t-1} - (1 - 4\eta\varepsilon + 4\eta^2\varepsilon^2)^{t-1}].$$

Note that as $\eta \rightarrow 0$ the variance tends to infinity. Next, we wish to bound the variance uniformly in t . This requires that $\varepsilon < 4\eta$, and in that case $\text{Var}(w_{\text{bad}}^t) \leq m_{\text{bad}}^2 (1 - 4\eta\varepsilon + \varepsilon^2)^t$. Applying Chebyshev's inequality, we get a bound on the magnitude of w_{bad}^t , namely $\Pr(w_{\text{bad}}^t > 2) \leq \text{Var}(w_{\text{bad}}^t)$, and running LBGA for $t = O(\log(1/s))$ rounds ensures this. Indeed, to bound $m_{\text{bad}}^2 (1 - 4\eta\varepsilon + \varepsilon^2)^t \leq s$ we choose

$$t \geq \frac{2 \log(m_{\text{bad}}) + \log(\frac{1}{s})}{\log(1 - 4\eta\varepsilon + \varepsilon^2)}.$$

Note that the denominator is nonnegative, and tends to zero as $\eta \rightarrow \varepsilon/4$. A similar analysis for w_{good}^t gives

$$\Pr(w_{\text{good}}^t < \mathbb{E}(w_{\text{good}}^t) - 1) \leq m_{\text{good}}^2 (1 + 4\eta\varepsilon + \varepsilon^2)^{t-1},$$

and ensuring this is at most s requires

$$t \geq \frac{2 \log(m_{\text{good}}) + \log(\frac{1}{s})}{\log(1 + 4\eta\varepsilon + \varepsilon^2)}.$$

Letting t be at least the maximum of these two quantities, we have that with probability at least $1 - s$,

$$p_{\text{bad}}^t = \frac{w_{\text{bad}}^t}{w_{\text{bad}}^t + w_{\text{good}}^t} \leq \frac{2}{m_{\text{good}}^2 (1 + 4\eta\varepsilon + \varepsilon^2)^t}.$$

And to ensure this is at most δ we need at least

$$t \geq \frac{1 + \log(1/\delta) - 2 \log(m_{\text{good}})}{\log(1 + 4\eta\varepsilon + \varepsilon^2)}.$$

So we pick t to be at least the maximum of all three of these lower bounds, and the proposition follows. \blacksquare

Corollary 1: Let $0 < \eta < 1/2$ and set $0 < \varepsilon < 4\eta$. Let n be the number of vertices of the input graphs, and suppose the number of input graphs m is $O(1)$. Suppose LBGA uses the quality function q_η , a null event A , and parameters $\nu = \varepsilon$. Suppose further that for every vertex pair (u, v) some input graph H_i is good for (u, v) . Then for any $0 < \delta < 1$, and any $0 < s < 1$, after $O(\log(n) + \log(1/s) + \log(1/\delta))$ rounds all u, v will satisfy $p_{u,v,\text{bad}} < \delta$ with probability at least $1 - s$.

Proof: Set $s' = s/n^2$. Run LBGA for $t = O(\log(1/s') + \log(1/\delta))$ rounds and apply the union bound over edges with Theorem 1. \blacksquare

One can further prove that for sparse noise, the neighborhood overlap function is a probabilistic oracle for the local stochastic block model. These theorems ignore the event A , and the end goal of this line of inquiry would characterize how a suitable event A can compensate for increasingly weak q and unreliable input graphs, since this is the real-world scenario for which we posit LBGA is useful. This is an area for future work.

VII. CONCLUSIONS

We present the Locally Boosted Graph Aggregation framework, a general framework for learning graph representations with respect to an application. In this paper, we demonstrate the strength of the framework with the application of community detection, and we believe the framework can be adapted to other inference goals in graphs such as link prediction or diffusion estimation.

Our framework offers a flexible, local weighting and aggregation of different edge sources in order to better represent the variability of relational structure observed in real networks. Inspired by concepts in boosting and bandit learning approaches, LBGA is designed to handle aggregations of noisy and disparate data sources, therefore marking a departure from methods that assume overlap and usefulness among all data sources considered.

LBGA also simplifies the task of designing a graph aggregation algorithm into designing a principled quality measure q and global event A . Doing so allows us to connect the utility of the graph representation to the application of interest. As a byproduct, we conjecture that statistics produced by LBGA provide information about the utility of the graph source, such as the level of noise. Such information can be used to improve the data collection process and partially mitigate the effects noise before it propagates to subsequent analysis. We gave evidence for the resistance of our framework to noise by running it on datasets with various known levels of noise, and observed the resulting matching weight distributions.

A primary direction for future work is to analyze the utility of our framework with respect to other graph applications, as well as to present a comparison of LBGA with other existing

multigraph clustering algorithms. Additional directions include a more thorough stability analysis of LBGA, exploring the modifications we have suggested to improve scalability, and to prove further theoretical results as described earlier.

VIII. ACKNOWLEDGEMENTS

We thank Vineet Mehta for providing us with the Enron topic model data, and for his many helpful discussions.

REFERENCES

- [1] Enron email dataset. <https://www.cs.cmu.edu/~enron/>.
- [2] C.C. Aggarwal, Y. Xie, and P.S. Yu. Towards community detection in locally heterogeneous networks. In *SDM*, pages 391–402. SIAM, 2011.
- [3] S. Arora, E. Hazan, and S. Kale. The multiplicative weights update method: a meta-algorithm and applications. *Theory of Computing*, 8(1):121–164, 2012.
- [4] M. Balcan and A. Blum. *On a theory of kernels as similarity functions*. Manuscript, 2006.
- [5] M. Balcan, A. Blum, and N. Srebro. A theory of learning with similarity functions. *Machine Learning*, 72(1-2):89–112, 2008.
- [6] M. Berlingerio, M. Coscia, and F. Giannotti. Finding redundant and complementary communities in multidimensional networks. In *CIKM*, pages 2181–2184, 2011.
- [7] S. Bubeck and N. Cesa-Bianchi. Regret analysis of stochastic and nonstochastic multi-armed bandit problems. *Foundations and Trends in Machine Learning*, 5(1):1–122, 2012.
- [8] S. Bubeck, R. Munos, and G. Stoltz. Pure exploration in multi-armed bandits problems. In *ALT*, pages 23–37, 2009.
- [9] R. S. Caceres, T. Y. Berger-Wolf, and R. Grossman. Temporal scale of processes in dynamic networks. In *ICDM Workshops*, pages 925–932, 2011.
- [10] D. Cai, Z. Shao, X. He, X. Yan, and Jiawei Han. Community mining from multi-relational networks. In *PKDD*, pages 445–452, 2005.
- [11] L. Danon, J. Duch, A. Diaz-Guilera, and A. Arenas. Comparing community structure identification. *J. Stat. Mech.*, 2005:P09008, 2005.
- [12] L. R. Dice. Measures of the amount of ecologic association between species. *Ecology*, 26(3):297–302, July 1945.
- [13] N. Eagle and A. Pentland. Reality mining: sensing complex social systems. *Personal and Ubiquitous Computing*, 10(4):255–268, 2006.
- [14] P. Erdős and A. Rényi. On random graphs, I. *Publicationes Mathematicae (Debrecen)*, 6:290–297, 1959.
- [15] S. Fortunato and M. Barthélemy. Resolution limit in community detection. *Proceedings of the National Academy of Sciences*, 104(1):36–41, 2007.
- [16] B. Gallagher, H. Tong, T. Eliassi-Rad, and C. Faloutsos. Using ghost edges for classification in sparsely labeled networks. In *Proc. (14th) ACM SIGKDD Inter. Conf. on Knowledge Discovery and Data Mining*, pages 256–264. ACM, 2008.
- [17] L. Getoor and C. P. Diehl. Link mining: a survey. *ACM SIGKDD Explorations Newsletter*, 7(2):3–12, 2005.
- [18] E. Gilbert and K. Karahalios. Predicting tie strength with social media. In *Proc. of SIGCHI Conf. on Human Factors in Computing Systems*, CHI ’09, pages 211–220, New York, NY, USA, 2009. ACM.
- [19] D. F. Gleich and C. Seshadhri. Vertex neighborhoods, low conductance cuts, and good seeds for local community methods. In *Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD ’12, pages 597–605, New York, NY, USA, 2012. ACM.
- [20] P. Jaccard. The distribution of the flora in the alpine zone. *New Phytologist*, 11(2):37–50, February 1912.
- [21] B. Klimt and Y. Yang. The enron corpus: A new dataset for email classification research. In *ECML*, pages 217–226, 2004.
- [22] T. G. Kolda and B. W. Bader. Tensor decompositions and applications. *SIAM review*, 51(3):455–500, 2009.
- [23] J. Leskovec, K. J. Lang, A. Dasgupta, and M. W. Mahoney. Statistical properties of community structure in large social and information networks. In *Proc. (17th) Inter. Conf. on World Wide Web*, WWW ’08, pages 695–704, New York, NY, USA, 2008. ACM.
- [24] M. Ley. The dblp computer science bibliography: Evolution, research issues, perspectives. In *SPIRE*, pages 1–10, 2002.
- [25] A. McCallum, A. Corrada-Emmanuel, and X. Wang. The author-recipient-topic model for topic and role discovery in social networks, with application to enron and academic email. In *Workshop on Link Analysis, Counterterrorism and Security*, pages 33–44, Newport Beach, CA, 2005.
- [26] A. K. McCallum. MALLET: A Machine Learning for Language Toolkit. <http://mallet.cs.umass.edu>, 2002.
- [27] B. A. Miller and N. Arcolano. Spectral subgraph detection with corrupt observations. In *Proc. IEEE Int. Conf. Acoust., Speech and Signal Process.*, 2014. To appear.
- [28] P. J. Mucha, T. Richardson, K. Macon, M. A. Porter, and J. P. Onnela. Community structure in time-dependent, multiscale, and multiplex networks. *Science*, 328(5980):876–878, 2010.
- [29] R. R. Nadakuditi and M. E. J. Newman. Graph spectra and the detectability of community structure in networks. *CoRR*, abs/1205.1813, 2012.
- [30] J. Neville and D. Jensen. Leveraging relational autocorrelation with latent group models. In *Proc. (4th) Inter. Workshop on Multi-relational Mining*, pages 49–55. ACM, 2005.
- [31] M. E. J. Newman. Modularity and community structure in networks. *Proceedings of the National Academy of Sciences of the United States of America*, 103(23):8577–8696, 2006.
- [32] E. E. Papalexakis, L. Akoglu, and D. Ience. Do more views of a graph help? community detection and clustering in multi-graphs. In *FUSION*, pages 899–905, 2013.
- [33] P. Pons and M. Latapy. Computing communities in large networks using random walks. *Journal of Graph Algorithms and Applications*, 10(2):191–218, 2006.
- [34] M. Rocklin and A. Pinar. On clustering on graphs with multiple edge types. *Internet Mathematics*, 9(1):82–112, 2013.
- [35] R. A. Rossi, L. McDowell, D. W. Aha, and J. Neville. Transforming graph data for statistical relational learning. *J. Artif. Intell. Res. (JAIR)*, 45:363–441, 2012.
- [36] R. E. Schapire. The strength of weak learnability. *Machine Learning*, 5:197–227, 1990.
- [37] L. Tang, X. Wang, and H. Liu. Community detection via heterogeneous interaction analysis. *Data Min. Knowl. Discov.*, 25(1):1–33, 2012.
- [38] W. Tang, Z. Lu, and I. S. Dhillon. Clustering with multiple graphs. In *Proc. (2009) IEEE Int. Conf. on Data Mining*, ICDM ’09, pages 1016–1021, Washington, DC, USA, 2009. IEEE Computer Society.
- [39] Y. Wang and G. Wong. Stochastic Block Models for Directed Graphs. *Journal of the American Statistical Association*, 82(397):8–19, 1987.
- [40] R. Xiang, J. Neville, and M. Rogati. Modeling relationship strength in online social networks. In *Proc. (19th) Int. Conf. on World Wide Web*, WWW ’10, pages 981–990, New York, NY, USA, 2010. ACM.