

Locally Boosted Graph Aggregation for Community Detection *

Jeremy Kun^{1,2}, Rajmonda S. Caceres¹, and Kevin M. Carter¹

¹MIT Lincoln Laboratory

²University of Illinois at Chicago

Abstract

Learning the right graph representation from noisy, multi-source data has garnered significant interest in recent years. A central tenet of this problem is relational learning. Here the objective is to incorporate the partial information each data source gives us in a way that captures the true underlying relationships. To address this challenge, we present a general, boosting-inspired framework for combining weak evidence of entity associations into a robust similarity metric. Building on previous work, we explore the extent to which different local quality measurements yield graph representations that are suitable for community detection. We present empirical results on a variety of datasets demonstrating the utility of this framework, especially with respect to real datasets where noise and scale present serious challenges. Finally, we prove a convergence theorem in an ideal setting and outline future research into other application domains.

1 Introduction

In the study of networks, the data used to define nodes and edges often come from multiple sources. These sources are often noisy and ambiguously useful, and the process of combining them into a single graph representation is critically important. For example, suppose we are studying a social network and wish to detect communities. The data that indicate membership in the same community are plentiful: communication, physical proximity, reported friendship, and many others. Each data source carries a different level of information about the underlying social structure, and each may accurately represent only some of the individuals. Some groups of friends communicate through Facebook and others via Instagram, etc. The best way to aggregate this information is far from clear, and recent work has shown that the choice of graph representation

*This work is sponsored by the Assistant Secretary of Defense for Research & Engineering under Air Force Contract FA8721-05-C-0002. Opinions, interpretations, conclusions and recommendations are those of the authors and are not necessarily endorsed by the United States Government.

significantly impacts the performance of subsequent machine learning and data mining algorithms [19, 18, 31, 11, 28].

Further complicating matters, the quality of the aggregated graph depends heavily on the application domain. Community detection is easiest with a graph representation that retains only inter-community edges, but to predict the spread of a virus one must have access to the strongest cross-community conduits. Suitable graph representations for these two tasks may come from the same data but are qualitatively different enough to warrant different aggregation techniques.

Even though the impact of the graph representation on subsequent analysis has been widely studied, there are few techniques for learning the right graph representations. Aggregation is often ad-hoc in practice, making it difficult even to compare algorithms within the same application domain using different data sources. The need for rigorous approaches to graph representation learning is even more apparent with big data, where variety and veracity compound the challenges of volume and velocity.

In this paper, we present a graph aggregation framework designed to make the process of learning a useful underlying graph representation rigorous with respect to application specific requirements. Our framework is called *Locally Boosted Graph Aggregation (LBGA)*. LBGA extracts the application-specific aspects of the learning objective as an event A representing an operation on the graph (e.g. a clustering algorithm, a random walk, etc.) and a local quality measure q . The framework incorporates this information into a reward system that promotes the presence of good edges and the absence of bad edges, in a fashion inspired by boosting.

Building on our work in [10], we demonstrate LBGA with the application of community detection. In this context the goal of graph representation learning is to aggregate the different data sources into a single graph which makes the true community structure easy to detect. LBGA evaluates the graph data locally, so that it can choose the data sources which most accurately represent the local structure of communities observed in real networks [3, 24]. In the absence of ground truth knowledge or one efficiently computable measure that can capture true community quality, LBGA relies on the pair of a graph clustering algorithm A and a local clustering metric q as an evaluation proxy. We show through empirical analysis that our algorithm can learn a high-quality global representation guided by the local quality measures considered.

We make the following contributions:

1. We present a graph aggregation framework that learns a useful graph representation with respect to an application requiring only a local heuristic measure of quality.
2. Our framework is stochastic and incorporates both edge and non-edge information, making it robust and suitable for sparse and noisy networks.
3. We demonstrate the success of an algorithm implementing the framework

for community detection, testing it against both synthetic data and real-world data sets.

4. We perform sensitivity and scalability analyses of our algorithm, showing that the algorithm scales linearly in the number of edges and is robust enough to handle large, noisy graphs.
5. We prove a convergence theorem for our framework and suggest the next steps in proving performance guarantees.

The rest of the paper is organized as follows. In Section 2 we give a brief overview of related literature. In Section 3 we discuss in detail the LBGA framework. In Section 4 we present the experimental analysis and results. In Section 5 we discuss sensitivity to noise and scalability. In Section 6 we prove our convergence theorem, and in Section 7 we discuss future work.

2 Related Work

2.1 Representation Learning and Clustering

Representation learning has garnered a lot of interest and research in recent years. Its goal is to introduce more rigor to the often ad-hoc practices of transforming raw, noisy, multi-source data into inputs for data mining and machine learning algorithms. Within this area, representation learning of graph-based data includes modeling decisions about the nodes of the graph, the edges, as well as the critical features that characterize them both.

In this context, Rossi et al. [36] discuss transformations to heterogeneous graphs (graphs with multiple node types and/or multiple edge types) in order to improve the quality of a learning algorithm. Within their taxonomy, our work falls under the link interpretation and link re-weighting algorithms [41, 20]. Our setting is different because we explicitly allow different edge types between the same pair of vertices. Also, our approach is stochastic, which we find necessary for learning a robust representation and weeding out noise.

Clustering in multi-edge graphs [33, 39, 38, 29, 7] is another area with close connections to our work. A common thread among these existing approaches is clustering by leveraging shared information across different graph representations of the same data. These approaches do not address scenarios where the information provided by the different sources is complementary or the overlap is scarce. In contrast, our approach iteratively selects those edge sources that lead to better clustering quality, independently of disagreement across the different features. [35, 12] present approaches for identifying the right graph aggregation, given a complete ground truth clustering, or a portion of it (i.e.: the cluster assignment is known only for a subset of the vertices in the graph). Our framework requires no such knowledge, but we do use ground truth to validate our experiments on synthetic data (Section 4.3).

Balcan and Blum define in [5, 6] a list of intuitive properties a similarity function needs to have in order to be able to cluster well. However, testing

whether a similarity function has the discussed properties is NP-hard, and often dependent on having ground truth available. Our model instead uses an efficiently computable heuristic as a rough guide.

2.2 Boosting and Bandits

Our framework departs from previous work most visibly through its algorithmic inspirations, namely boosting [37] and bandit learning (see [8] for a survey of the latter). In boosting, we assume we have a *weak classifier*, whose performance is only slightly better than random. In his landmark paper [37], Schapire proved that weak classifiers can be combined via a majority voting scheme to form a learner that achieves arbitrarily close to perfect generalization. We think of different graph data sources as weak learners in that they offer knowledge on when an edge should be present. Then the question becomes whether one can “boost” the knowledge in the different graphs to make one graph representation that is arbitrarily good.

Unfortunately, our problem setting does not allow pure boosting for two reasons. First, boosting assumes the learners are all slightly better than random, but graph representations can be pure noise or can even provide *bad* advice. Second, boosting has access to ground truth classification labels. Even with reliable input data, the quality of the aggregation depends on the application and many applications have no standard measure of quality.

Ideas from bandit learning compensate for these problems. In bandit learning an algorithm receives rewards as it explores a set of actions, and the goal is to minimize some notion of regret in hindsight. The basic model has many variants, but two central extensions in the literature are expert advice and adversaries. Expert advice consists of functions suggesting to the algorithm what action to take in each round (e.g., weak classifiers). The adversarial setting involves an adversary who knows everything but the random choices made by the algorithm in advance, and sets the experts or rewards so as to incur maximum regret.

We apply these ideas to graph representation learning by setting up an artificial reward system based on the given application and using stochasticity to weed out adversarial portions of the input graphs. In our setting we only care if the graph representation is good at the end, while bandit learning often seeks to maximize cumulative rewards during learning. There are bandit settings that only care about the final result (e.g., the pure exploration model of Bubeck et al. [9]), but to the best of our knowledge no theoretical results in the bandit literature immediately apply to our framework. This is largely because we rely on heuristic proxies to measure the quality of a graph, so even if the bandit learning objective is optimized we cannot guarantee the result is useful.¹ Nevertheless we can adapt the successful techniques and algorithms for boosting and bandit learning, and hope they produce useful graphs in practice. As the rest of this paper demonstrates, they do indeed.

¹For example, the empty graph maximizes some proxies but is entirely useless.

The primary technique we adapt from bandits and boosting is the Multiplicative Weights Update Algorithm (MWUA) [4]. The algorithm works as follows. A list of weights is maintained on each element x_j of a finite set X . At each step of some process an element x_i is chosen (in our case, by normalizing the weights to a probability distribution and sampling), a reward $q_{t,i}$ is received, and the weight for x_i is multiplied or divided by $(1 + \varepsilon q_{t,i})$, where $\varepsilon > 0$ is a fixed parameter controlling the rate of update. After many rounds, the elements with the highest weight are deemed the best and used for whatever purpose needed.

3 The Locally Boosted Graph Aggregation Framework

The Locally Boosted Graph Aggregation framework (LBGA) can succinctly be described as running MWUA for each possible edge, forming a candidate graph representation G_t in each round by sampling from all edge distributions, and computing local rewards on G_t to update the weights for the next round. Over time G_t stabilizes and we produce it as output. The remainder of this section expands the details of this sketch and our specific algorithm implementing it.

3.1 Framework Details

Let H_1, \dots, H_m be a set of unweighted², undirected graphs defined on the same vertex set V . We think of each H_i as “expert advice” suggesting for any pair of vertices $u, v \in V$ whether to include edge $e = (u, v)$ or not. Our primary goal is to combine the information present in the H_i to produce a global graph representation G^* suitable for a given application.

We present LBGA in the context of community detection, noting what aspects can be generalized. Each round has four parts: producing the aggregate candidate graph G_t , computing a clustering A for use in measuring the quality of G_t , computing the local quality of each edge, and using the quality values to update the weights for the edges. After some number of rounds T , the process ends and we produce $G^* = G_T$.

Aggregated Candidate Graph G_t : In each round produce a graph G_t as follows. Maintain a non-negative weight $w_{u,v,i}$ for each graph H_i and each edge (u, v) in $H_1 \cup \dots \cup H_m$. Normalize the set of all weights for an edge $\mathbf{w}_{u,v}$ to a probability distribution over the H_i ; thus one can sample an H_i proportionally to its weight. For each edge, sample in this way and include the edge in G_t if it is present in the drawn H_i .

Event $A(G_t)$: After the graph G_t is produced, run a clustering algorithm A on it to produce a clustering $A(G_t)$. In this paper we fix A to be the Walktrap algorithm [34], though we have observed the effectiveness of other clustering algorithms as well. In general A can be any event, and in this case we tie it to the application by making it a simple clustering algorithm.

²There is a natural extension for weighted graphs.

Local quality measure: Define a *local quality measure* $q(G, e, c)$ to be a $[0, 1]$ -valued function of a graph G , an edge e of G , and a clustering c of the vertices of G . The quality of (u, v) in G_t is the “reward” for that edge, and it is used to update the weights of each input graph H_i . More precisely, the reward for (u, v) in round t is $q(G_t, (u, v), A(G_t))$.

Update Rule: Update the weights using MWUA as follows. Define two learning rate parameters $\varepsilon > 0, \nu > 0$, with the former being used to update edges from G_t that are present in H_i and the latter for edges not in H_i . In particular, suppose $q_{u,v}$ is the quality of the edge (u, v) in G_t . Then, the update rule is defined as follows:

$$w_{u,v,i} = \begin{cases} w_{u,v,i}(1 + \varepsilon q_{u,v}), & \text{if } (u, v) \in H_i \\ w_{u,v,i}(1 - \nu q_{u,v}), & \text{if } (u, v) \notin H_i. \end{cases}$$

3.2 Quality Measures for Community Detection

We presently describe the two local quality measures we use for community detection. The first, which we call *Edge Consistency (EC)* captures the notion that edges with endpoints in the same cluster are superior to edges across clusters:

$$EC_{u,v} = \begin{cases} 1, & \text{if } c(u) = c(v) \\ 0, & \text{if } c(u) \neq c(v). \end{cases}$$

EC offers a quality metric that is inextricably tied to the performance of the chosen clustering algorithm. The idea behind edge consistency can also be combined with any quality function q to produce a “consistent” version of q . Simply evaluate q when the edge is within a cluster, and $-q$ when the edge is across clusters. Note that q need not depend on a clustering of the graph or the clustering algorithm, and it can represent algorithmic-agnostic measures of clustering quality.

As an example of such a measure q , we consider the metric of *Neighborhood Overlap (NO)*, which uses the idea that vertices that share many neighbors are likely to be in the same community. NO declares that the quality of (u, v) is equal to the (normalized) cardinality of the intersection of the neighborhoods of u and v :

$$NO_{u,v} = \frac{|N(u) \cap N(v)|}{|N(u) \cap N(v)| + \log(|V|)},$$

where $N(x)$ represents the neighborhood of vertex x . We have also run experiments using more conventional normalizing mechanisms, such as the Dice and Jaccard indices [14, 22]), but our neighborhood overlap metric outperforms them by at least 10% in our experiments. We argue this is due to the use of a global normalization factor, as opposed to a local one, which is what Dice and Jaccard indices use. For brevity and simplicity, we omit our results for Jaccard and Dice indices and focus on Neighborhood Overlap. In our experimental

analysis (Section 4.4) we use the consistent version of *NO*, which we denote *consistentNO*.

While we demonstrate the utility of the LBGA framework by using *EC* and *consistentNO*, the design of the framework is modular, in that the mechanism for rewarding the “right” edges is independent from the definition of reward. This allows us to plug in other quality metrics to guide the graph representation learning process for other applications, a key goal in LBGA’s design.

3.3 LBGA Implementation

Processing every edge in every round of the LBGA framework is inefficient. Our implementation of LBGA, given by Algorithm 1, improves efficiency by fixing edges whose weights have grown so extreme so as to be picked with overwhelming or negligible probability (with probability $> 1 - \delta$ or $< \delta$ for a new parameter δ). In practice this produces a dramatic speedup on the total runtime of the algorithm.³ The worst-case time complexity is the same, but balancing parallelization and the learning parameters suffices for practical applications.

In addition, our decision to penalize non-edges ($\nu > 0$) also improves runtime from the alternative ($\nu = 0$). In our experiments non-edge feedback causes G_t to convergence in roughly half as many rounds as when only presence of edge is considered as indication of relational structure.

We also note that Algorithm 1 stays inside the “boundaries” determined by the input graphs H_i . It never considers edges that are not suggested by *some* H_i , nor does it reject an edge suggest by all H_i . Thus, when we discuss sparsity of our algorithm’s output in our experiments, we mean with respect to the number of edges in the union of the input graphs.

4 Experimental Analysis

We presently describe the datasets used for analysis and provide quantitative results for the performance of Algorithm 1.

4.1 Synthetic Datasets

Our primary synthetic data model is the stochastic block model [40], commonly used to model explicit community structure. We construct a probability distribution $G(\mathbf{n}, B)$ over graphs as follows. Given a number n of vertices and a list of cluster (block) sizes $\mathbf{n} = \{n_1, \dots, n_k\}$ such that $n = \sum_i n_i$, we partition the n vertices into k blocks $\{b_1, \dots, b_k\}$, $|b_i| = n_i$. We declare that the probability of an edge occurring between a vertex in block b_i and block b_j is given by the (i, j) entry of a k -by- k matrix B . In order to simulate different scenarios, we consider the following three cases.

³From days to minutes in our experiments.

```

Data: Unweighted graphs  $H_1, \dots, H_m$  on the same vertex set  $V$ , a clustering
algorithm  $A$ , a local quality metric  $q$ , three parameters  $\varepsilon, \nu, \delta > 0$ 
Result: A graph  $G$ 
Initialize a vector  $\mathbf{w}_{u,v} = \mathbf{1}$  for all  $u \neq v \in V$ 
Let  $U$  be the edge set of  $H_1 \cup \dots \cup H_m$ 
Let  $G_{\text{learned}} = (V, \emptyset)$ 
while  $|U| > 0$  do
    Let  $G$  be a copy of  $G_{\text{learned}}$ 
    for  $(u, v) \in U$  do
        Let  $p_{u,v} = \frac{\sum_i w_{u,v,i} 1_{\{(u,v) \in H_i\}}}{\sum_i w_{u,v,i}}$ 
        Flip a coin with bias  $p_{u,v}$ 
        If heads, include  $(u, v)$  in  $G$ .
    end
    Cluster  $G$  using  $A$ 
    for  $(u, v) \in U$  do
        Set  $p = q(G, A(G), (u, v))$ 
        for  $i = 1, \dots, m$  do
            if  $(u, v) \in H_i$  then
                Set  $w_{u,v,i} = w_{u,v,i}(1 + \varepsilon p)$ 
            else
                Set  $w_{u,v,i} = w_{u,v,i}(1 - \nu p)$ 
            end
        end
        Let  $p_{u,v} = \frac{\sum_i w_{u,v,i} 1_{\{(u,v) \in H_i\}}}{\sum_i w_{u,v,i}}$ 
        if  $p_{u,v} > 1 - \delta$  then
            Add  $(u, v)$  to  $G_{\text{learned}}$ , remove it from  $U$ 
        end
        if  $p_{u,v} < \delta$  then
            Remove  $(u, v)$  from  $U$ 
        end
    end
end
Output  $G$ 

```

Algorithm 1: Optimized implementation of LBGA. Note that 1_E denotes the characteristic function of the event E .

Global Stochastic Block Model (GSBM): In this model we have m input graphs H_1, \dots, H_m , each drawn from the stochastic block model $G(\mathbf{n}, B_i)$ ⁴, with $n_1 = \dots = n_m$ and B_i defined as:

$$B_i = \begin{pmatrix} p_i & r_i & r_i & \dots & r_i \\ r_i & p_i & r_i & \dots & r_i \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ r_i & r_i & r_i & \dots & p_i \end{pmatrix},$$

where p_i represents the within-cluster edge probability and r_i represents the across-cluster edge probability in graph H_i . The ratio $SNR = p_i/r_i$ is commonly referred to as the *signal to noise* ratio and captures the strength of community

⁴ $G(\mathbf{n}, B_i)$ represents a simpler case of the stochastic block model, where the within-cluster probabilities are uniform across blocks and blocks have the same size.

Dataset	Parameters
GSBM-1	$m = k = 4, n_i = 125, p_i = 0.2, r_i = 0.05, i = 1, \dots, m$
GSBM-2	$m = k = 4, n_i = 125, p_i = 0.3, r_i = 0.05, i = 1, \dots, m$
GSBM-3	$m = 5, k = 4, n_i = 125, p_i = 0.3, r_i = 0.05, i = 1, \dots, 4, p_5 = r_5 = 0.01$
GSBM-4	$m = k = 4, n_i = 125, p_1 = 0.1625, p_2 = 0.125, p_3 = 0.125, p_4 = 0.0875, r_i = 0.05, i = 1, \dots, m$
GSBM-5	$m = k = 4, n_i = 125, p_1 = 0.15, p_2 = 0.1, p_3 = p_4 = 0.05, r_i = 0.05, i = 1, \dots, m$
LSBM-1	$m = k = 4, n_i = 125, p_i = 0.2, r_i = 0.05, i = 1, \dots, m$
LSBM-2	$m = k = 4, n_i = 125, p_i = 0.3, r_i = 0.05, i = 1, \dots, m$
LSBM-3	$m = 5, k = 4, n_i = 125, p_i = 0.3, r_i = 0.05, i = 1, \dots, m, p_5 = r_5 = 0.01$
ER only	$m = 4, p_i = r_i = 0.01$
DBLP	$n = 3153, m = 2$
RealityMining	$n = 90, m = 6$
Enron	$n = 145, m = 2, \alpha = 0.9$

Table 1: Description of datasets analyzed. Total number of vertices in each synthetic source graph is $n = 500$. m is the number of graph sources. k is the number of clusters. n_i represents number of vertices in cluster i . p_i and r_i represent the within- and across-cluster edge probability for each the m graph sources.

structure within H_i . We use the GSBM case to model a scenario where each graph source has a global (or uniform) contribution toward the quality of the targeted graph representation G^* .

Local Stochastic Block Model (LSBM): This scenario captures the notion that one graph source accurately describes one community, while another source fares better for a different community. For example, if we have two underlying communities, and two graph sources H_1, H_2 , then we use the following two block matrices to represent them:

$$B_1 = \begin{pmatrix} p & r \\ r & r \end{pmatrix}, \quad B_2 = \begin{pmatrix} r & r \\ r & p \end{pmatrix}.$$

This naturally extends to a general formulation of the LSBM model for m communities.

Erdős-Rényi (ER) model: Finally, we consider the case of the Erdős-Rényi random graph [16], where any two vertices have equal probability of being connected. This model provides an example of a graph with no community structure. Note that the ER model is a special case of both GSBM and LSBM with $p = r$. In our experimental analysis we consider cases where an ER model is injected into instances of GSBM and LSBM in order to capture a range of structure and noise combinations.

4.2 Real Datasets

4.2.1 DBLP

Our first real-world dataset is DBLP [25], a comprehensive online database documenting research in computer science. We extracted the subset of the DBLP database corresponding to researchers who have published at two conferences: the Symposium on the Theory of Computing (STOC), and the Symposium on Foundations of Computer Science (FOCS). The breadth of topics presented at these conferences implies a natural community structure organized by sub-field. Each node in the DBLP graph represents an author, and we use two graphs on this vertex set: the *co-authorship* graph and the *title similarity* graph. For the

latter, we add an edge between two author vertices if any of their paper titles contain at least three words in common (excluding stop words), and the weight of this edge is the number of such pairs of papers. We considered a total of 5234 papers across 3153 researchers.

4.2.2 RealityMining

Our second dataset is RealityMining [15], a 9-month experiment in 2004 which tracked a group of 90 individuals at MIT via sensors in their cell phones. The individuals were either associated with the MIT Media Lab or the Sloan Business School, and there is a natural corresponding community structure. The data collected include voice calls, bluetooth scan events at five-minute intervals, cell tower usage, and self-reported friendship and proximity data. The data set is naturally noisy: surveys are subjective estimates, cell tower ranges are only so precise and signal outages are common, and there was data loss from typical cell phone problems like running out of battery.

We used the subset of subjects participating between 2004-09-23 19:00:00 and 2005-01-07 18:00:00 (UTC-05:00), for a total of 3354 call events, 786301 cell tower transition events, and 689025 bluetooth scan events. The nodes in our graphs represent individuals in the study. Weighted edges correspond to the total duration of voice calls, the total amount of time two individuals used the same cell tower, the total number of bluetooth events, and the results of the friendship/proximity surveys for a total of six graphs. Our results were stable under somewhat drastic changes in the input graphs (for example, in ignoring weights).

4.2.3 Enron

Our final dataset is the Enron email dataset [23, 1], a well-studied corpus of over 600,000 emails sent between 145 employees of the Enron Corporation in the early 2000's. We produced two graphs from the Enron data, one for peer-to-peer email communication and one for topic similarity in the email content.

In both graphs the vertices are individuals. In the email graph, the edges are weighted by the number of emails sent between the individuals in question. We used the Mallet package [27] to generate the LDA topic model for the content of Enron email data. We aggregated into one document all the email content sent by each of the Enron employees considered in the email link graph. Each document and therefore each sender is represented by 60 topics. We measure cosine distance of the topic vectors of individuals, and considered an edge as present if the cosine distance was above a specified threshold value α .⁵

Table 1 contains a summary of all the datasets used for the experimental analysis and their parameters.

⁵We experimented with various threshold values, and we discuss this in Section 4.4.4.

4.3 Validation Procedure

In our work, the optimality of the graph representation is closely coupled with the quality of community structure captured by the representation. This gives us several ways of evaluating the quality of the results produced by our algorithm. We consider notions of quality reflected at different levels: the quality of cluster assignment, the quality of graph representation, and the quality of graph source weighting.

Quality of Cluster Assignment: Since the output of LBGA is a graph, we use the walktrap clustering algorithm to extract communities for analysis. We then compare these communities to the ground truth clustering, when it is available, or else to the known features of the datasets. We use the Normalized Mutual Information (NMI) measure [13] to capture how well the ground truth clustering overlaps with the clustering on the graph representation output from our algorithm. In general we find that the choice of clustering algorithm is unimportant, because the graph output by LBGA is sufficiently modular to admit only one reasonable clustering. Walktrap is further convenient in that we need not assume the number of clusters ahead of time.

Quality of Graph Representation: An ideal graph representation that contains community structure would consist of disjoint cliques or near-cliques corresponding to the communities. As we illustrate in Section 4.4, an optimal graph representation can do better than just produce a perfect clustering. It can also remove cross-community edges and produce a sparser representation, which is what our algorithm does. We use two measures of clusterness to capture this notion of graph representation quality. Modularity [32] is a popular measure that compares a given graph and clustering to a null model. Conductance [24, 21] measures relative sparsity of cuts in the graph. Since conductance is defined for a single cut, we compute it for a clustering as the sum of the conductance values of cuts defined by isolating a single cluster from the rest of the graph. Note that *higher* modularity scores and *lower* conductance scores signify stronger community structure. Both modularity and conductance are well-known and often offer complimentary information about the quality of communities.

We note two extreme graph representation cases, the empty graph which is perfectly modular in a degenerate sense, and the union graph which is a trivial aggregation. To signal these cases in our results, we display the *sparsity* of the produced graph G^* , defined as the fraction of edges in G^* out of the total set of edges in all input graphs.

Quality of Graph Source Weighting: the quality of the aggregation process is captured by the right weighting of individual edge sources. Edge sources (input graphs) that are more influential in uncovering the underlying community structure have higher weights on average. Similarly, edge types that contribute equally should have equal weights, and edge types with no underlying structure should have low weights.




Figure 1: Left: the results of LBGA on the RealityMining dataset. Right: the input graph of Bluetooth scan events. LBGA was run with *consistentNO*, $\nu = \varepsilon = 0.2$, $\delta = 0.05$




Figure 2: Left: the results of LBGA on the Enron dataset. Right: the input graph of topic models thresholded at 0.9. LBGA was run with *consistentNO*, $\nu = \varepsilon = 0.2$, $\delta = 0.05$

4.4 Experimental Results

Table 2 contains the numerical results of our experiments. As a baseline, we computed the modularity values of the union of the input graphs with respect to the ground truth (for synthetic) or the Walktrap clusterings (for the real world). We now discuss the specific results for the synthetic and real data sets.

4.4.1 Synthetic

For illustration, we show in Figure 3 the performance of Algorithm 1 when *consistentNO* is used as a local quality metric and LSBM-3 (see Table 1 for details) is used to generate the input graphs. Note that the algorithm converges quickly to a graph which results in a perfect clustering as measured by NMI. We also plot the modularity of the resulting graph produced in each round, seeing that it far exceeds the “baseline” modularity of the union of the input graphs. This tells us the learning algorithm is able to discard the noisy edges in the model. Finally, we plot the number of edges in the graph produced in each round, and the average vertex-pair weight for each input graph. This verifies that our algorithm complies with our edge-type weighting and sparsity requirements. Indeed, the algorithm produces a relatively sparse graph, using about 40% of the total edges available and weights edges from the Erdős-Rényi source

appropriately. Our algorithm hence achieves a superior graph than the union, while preserving the underlying community structure so as to be amenable to clustering.

The results for the other synthetic datasets are similar and summarized in Table 2. We note that our algorithm’s performance degrades when the noise becomes too high. In Section 5.1 we analyze the signal to noise ratio in the synthetic data sets more closely.

4.4.2 DBLP

In Figure 4, we show results for the DBLP dataset. For both edge consistency and *consistentNO*, our algorithm converges to a graph of modularity exceeding that of the union graph and using significantly fewer edges (60% in the case of *consistentNO* and 88% for edge consistency).

Our algorithm selects title similarity as having more influence in recovering communities for the STOC/FOCS conferences. Researchers attending these conferences represent a small community as a whole with many of them sharing co-authorship on papers with diverse topics. In this sense, it is not surprising that title similarity serves as a better proxy for capturing the more pronounced division along topics. We have also manually inspected the resulting clusters, and most appear organized both by membership and coauthorship. Take, for example, Mikko Koivisto, Thore Husfeldt, Petteri Kaski, and Andreas Björklund. They have together coauthored over 15 papers in combinatorial optimization, and naturally fall within the same small coauthorship cluster. However, the title similarity graph alone yields around 1500 clusters, and these researchers are split across two clusters because of the differences in their non-coauthored work. They fall in the same cluster in the LBGA-aggregated graph, and the cluster is larger, including well-known researchers who are either coauthors with some of the four or have done much work in the same field. Though this is a promising sign, we suggest a more thorough quantitative analysis of the quality of this clustering for future work.

4.4.3 RealityMining

Much work has been done in manually constructing good graph representations for the RealityMining social network (e.g. aggregating Bluetooth via thresholding and picking useful time windows [15, 11]). LBGA, however, arrives at an equally good representation from the raw, noisy input graphs. The final graph it constructs contains two dense clusters corresponding exactly to the MIT Media Lab and the Sloan Business School, with only three edges crossing the cut. The final group uses 63.5% of the total edges available and has modularity 0.25 with respect to the Media/Sloan partition. See Figure 1 for a comparison of the original Bluetooth graph and the final produced graph.

4.4.4 Enron

The data of Table 2 shows that *consistentNO* achieves a graph representation with higher modularity, lower conductance and better sparsity when compared to the edge consistency measure. Figure 2 shows a clear community structure, and there the smaller clusters correspond to lower-level employees while the higher level managers reside in the bigger clusters. Additionally, there was a known fantasy sports community within the network, and all of these individuals fall within a single cluster in the graph output by LBGA [26].

We also investigated the effect of changing the threshold value α for considering an edge in the topic graph. A value of less than $\alpha = 0.7$ always produced two large dense clusters with many noisy edges between them (akin to the two large clusters in Figure 2). In our experiments we used $\alpha = 0.9$, although values of α as high as 0.95 gave qualitatively similar results.

4.4.5 Comments

Overall, we find that LBGA converges to graphs of both high modularity and low conductance score. It also generates graph representations that induce correct clusterings in almost all cases where some sort of ground truth is known, the challenging case being when SNR is low. Moreover the algorithm weights the different input graphs appropriately to their usefulness. We find that the edge consistency measure outperforms neighborhood overlap in terms of overlap with ground truth clustering when the signal to noise ratio is particularly low, but that in almost all other cases (especially the real data sets), *consistentNO* produces higher quality, sparser, and more modular graphs.

We notice that the modularity values for RealityMining and the Enron data set are significantly smaller after passing through LBGA when compared to the baseline union values. We argue that this is due to the relatively small clusters produced by LBGA, as it is a known shortcoming that modularity is not an accurate measure when the communities are small [17]. Indeed, when conductance is used as a clustering quality measure LBGA significantly outperforms the baseline union aggregation. Figure 1 shows the favorable structure of the RealityMining dataset after being passed through LBGA (and the noisy input data), and Figure 2 gives a similar picture for the Enron dataset.

5 Robustness and Scalability

5.1 Sensitivity Analysis

We analyze the sensitivity of LBGA to noise. In Figure 5 we display performance as measured by NMI when the graph inputs are LSBM models across different intra-cluster probability values p_i and varying SNR values. We make the following general observations. The algorithm is consistent in that as the noise rate r_i increases, the NMI values do deteriorate as expected. The algorithm both reaches higher quality and maintains the quality longer for denser




Figure 3: Graph representation learning for LSBM-3. The LBGA parameters are $\varepsilon = \nu = 0.2, \delta = 0.05$. Plots in order top to bottom: 1. NMI of $A(G_t)$ with the ground truth clustering, 2. modularity of G_t w.r.t $A(G_t)$, with the horizontal line showing the modularity of the union of the input graphs w.r.t. ground truth, 3. the number of edges in G_t , 4. the average probability weight (quality) of vertex pairs for H_i . The Erdős-Rényi graph converges to low weight by round 300.

graphs, which again is consistent with our expectations. It seems the critical SNR (when the NMI drastically drops) is higher for consistentNO than for EC. At a signal to noise ratio of 2 or less, the NMI is bad for both quality measures and all choices of p_i . The sharp drop in quality is related to the well-known phase transition in the community detection problem [30]. Therefore, the distance from the theoretical detectability bound on community detection could be used as an agnostic measure to gauge the usefulness of LBGA with a particular quality metric.

5.2 Scalability Analysis

We run a larger version of the LSBM-2 model ($m = k = 10, p_i = 0.3, r_i = 0.05$) to illustrate how LBGA scales for larger graphs. As shown in Figure 6, LBGA scales linearly with the number of edges. Given that real-world graphs are usually sparse, this scaling behavior makes LBGA computationally suitable for large graphs.




Figure 4: Aggregation of co-authorship (red curve) and title similarity graphs (green curve) for DBLP dataset.

LBGA was designed and implemented with scaling in mind, and the process of fixing vertex-pairs as they are learned is largely the reason for the nice scaling properties of Algorithm 1. Moreover, the parameter δ encapsulates a trade-off between runtime and accuracy. Should linear scaling be insufficient, LBGA’s design allows for additional modifications to improve scalability. For example: quality functions that are sufficiently local allow one to parallelize the weight update step; one could sample a sublinear number of edges in each round and only update weights within the subgraph; one might specify a set of “seeds” (vertices of interest), and sample edges local to those vertices. While these methods are currently just potential directions for future work, there is clearly much potential for further improvements of LBGA’s computational efficiency.

6 A convergence theorem

In this section we provide a theorem on the convergence rate of LBGA in the case that the quality function actually has access to the ground truth clustering. This is an admittedly unrealistic scenario, but it serves a few purposes: it is the simplest possible theorem one might hope to prove, it is a demonstration of how standard MWUA analysis can be used for our problem, and it makes the

Dataset	Union Graph		EC					ConsistentNO				
	Modularity	Conductance	Modularity	Conductance	NMI	Sparsity	Edge Type Weights	Modularity	Conductance	NMI	Sparsity	Edge Type Weights
GSBM-1	0.264	7.568	0.549	2.042	1	0.844	(0.250,0.251,0.250,0.249)	0.750	0.002	1	0.515	(0.250,0.251,0.249,0.249)
GSBM-2	0.323	5.887	0.580	1.586	1	0.691	(0.252,0.250,0.248,0.251)	0.750	0	1	0.573	(0.252,0.250,0.247,0.251)
GSBM-3	0.312	6.140	0.607	1.333	1	0.657	(0.225,0.224,0.226,0.227,0.098)	0.750	0	1	0.562	(0.221,0.221,0.222,0.223,0.113)
GSBM-4	0.143	12.712	0.421	2.659	0.966	0.585	(0.302,0.237,0.305,0.302)	0.750	0.001	0.983	0.393	(0.302,0.231,0.306,0.302)
GSBM-5	0.145	6.584	0.395	2.124	0.919	0.653	(0.213,0.282,0.361,0.144)	0.666	0.001	0.958	0.477	(0.199,0.271,0.348,0.182)
LSBM-1	0.111	14.288	0.298	4.316	0.765	0.651	(0.253,0.250,0.250,0.248)	0.378	122.601	0.032	0.069	(0.249,0.251,0.250,0.250)
LSBM-2	0.167	11.106	0.464	3.290	0.975	0.582	(0.249,0.211,0.248,0.222)	0.750	0.001	1	0.417	(0.250,0.250,0.248,0.252)
LSBM-3	0.162	11.701	0.473	3.009	0.966	0.568	(0.218,0.217,0.222,0.219,0.124)	0.750	0.001	0.968	0.395	(0.212,0.212,0.213,0.209,0.154)
ER only	-0.002	24.729	0.193	112.947	0.012	0.999	(0.264,0.234,0.260,0.243)	0.836	1.068	0.025	0.230	(0.251,0.253,0.248,0.247)
DBLP	0.386	1368.839	0.372	1214.824	NA	0.962	(0.270,0.730)	0.695	159.286	NA	0.632	(0.432,0.568)
RealityMining	0.452	70.314	0.196	1.538	NA	0.724	(0.394,0.080,0.226,0.100,0.100,0.100)	0.246	0	NA	0.646	(0.365,0.091,0.198,0.115,0.115,0.115)
Enron	0.559	134.572	0.190	11.092	NA	0.921	(0.193,0.807)	0.444	0.594	NA	0.631	(0.390,0.610)

Table 2: LBGA performance results. All datasets in this table were run with EC and *consistentNO* using $\varepsilon = \nu = 0.2, \delta = 0.05$. Union modularity and conductance for real datasets was computed with the walktrap clustering. The order of edge type weights for the real datasets are: DBLP (coauthorship, title similarity); RealityMining (bluetooth, phone calls, cell tower proximity, reported friendship, in-lab proximity, out-lab proximity); Enron (email, topic similarity).

theoretical goals and assumptions rigorous.

Let H_1, \dots, H_N be graphs on the same vertex set V , and c an unknown target clustering of V . Call H_i *good* for a vertex pair (u, v) if it agrees with c , i.e. (u, v) is both an edge of H_i and $c(u) = c(v)$, or (u, v) is not an edge of H_i and $c(u) \neq c(v)$. Otherwise call H_i *bad* for (u, v) .

Let $\delta > 0$ be the learning target, so that an edge is considered “learned” if the probability p_i^t of choosing it in round t either drops below δ or rises above $1 - \delta$. Define the oracle-type quality function

$$q_c(u, v) = \begin{cases} 1 & \text{if } c(u) = c(v) \\ -1 & \text{if } c(u) \neq c(v) \end{cases}$$

The theorem is that with access to q_c LBGA will produce the correct graph after $O(\log(1/\delta))$ rounds. Before we prove this we set up a bit of notation.

Fix a vertex pair (u, v) . Let $w_{u,v,i}$ be the weight of H_i that LBGA maintains. Define $w_{u,v,\text{good}} = \sum_{\text{good } H_i} w_i$ to be the sum of the weights of the good graphs for (u, v) , and $w_{u,v,\text{bad}} = \sum_{\text{bad } H_i} w_i$ the sum of the bad weights. Call $p_{u,v,\text{good}}, p_{u,v,\text{bad}}$ the corresponding probability weights of picking a good or bad input graph. When u, v are clear from the context, we will omit them from the subscripts and simply write w_{good} , etc.

Theorem 1. *Suppose LBGA uses the quality function q_c , a null event A , and parameters $0 < \varepsilon = \nu < 1/2$. Suppose further that for every vertex pair (u, v) some input graph H_i is good for (u, v) . Then for any $0 < \delta < 1$ and for all vertex pairs (u, v) , after $O(\log(1/\delta))$ rounds the probability $p_{u,v,\text{bad}} < \delta$.*

Proof. Fix a vertex pair (u, v) . We add a superscript t to any quantity that changes over rounds to denote which round of LBGA the quantity refers to. Our goal is to bound p_{bad}^t for $t = O(\log(1/\delta))$. Denote by $N_{\text{bad}} = w_{\text{bad}}^1, N_{\text{good}} = w_{\text{good}}^1$, the number of bad and good input graphs, respectively.

The update step in LBGA and the quality function q_c are such that in every round the weights of good graphs are multiplied by $(1 + \varepsilon)$ and the weights of the bad graphs are multiplied by $(1 - \varepsilon)$. Hence,




Figure 5: Performance of LBGA (measured by NMI) as a function of SNR for the LSBM model with different probabilities p_i for the edge consistency measure (top) and *consistentNO* (bottom).

$$w_{\text{good}}^{t+1} = \sum_{H_i \text{ is good}} w_i^t (1 + \varepsilon) = w_{\text{good}}^t (1 + \varepsilon)$$

and likewise for w_{bad}^{t+1} . By induction, we have

$$w_{\text{good}}^t = N_{\text{good}} (1 + \varepsilon)^t, w_{\text{bad}}^t = N_{\text{bad}} (1 - \varepsilon)^t$$

and the value

$$p_{\text{bad}}^t = \frac{N_{\text{bad}} (1 - \varepsilon)^t}{N_{\text{bad}} (1 - \varepsilon)^t + N_{\text{good}} (1 + \varepsilon)^t} \leq \frac{N_{\text{bad}}}{(1 + \varepsilon)^t}$$

holds for all $t, N_{\text{good}}, N_{\text{bad}} \in \mathbb{N}, \varepsilon > 0$. If we fix a stopping point $T > 0$ and require $N_{\text{bad}} / (1 + \varepsilon)^T \leq \delta$ then we solve for

$$T \geq \frac{\log(1/\delta) + \log(N_{\text{bad}})}{\log(1 + \varepsilon)} = O(\log(1/\delta))$$

□

Note that this result is deterministic. The oracle-type quality function allows us to avoid a probabilistic analysis. So the natural generalization is to weaken




Figure 6: The runtime of LBGA as a function of number of edges on the union of input graphs.

the assumption on q . For example, one might assume its value correlates with the truth instead of being absolutely correct, or that its usefulness is determined i.i.d. in each round. The end goal of this line of inquiry would characterize how a suitable event A can compensate for increasingly weak and unreliable q , since this is the real-world scenario for which we posit LBGA is useful. This is an area for future work.

7 Conclusions

We present the Locally Boosted Graph Aggregation framework, a general framework for learning graph representations with respect to an application. In this paper, we demonstrate the strength of the framework with the application of community detection, and we believe the framework can be adapted to other inference goals in graphs such as link prediction or diffusion estimation.

Our framework offers a flexible, local weighting and aggregation of different edge sources in order to better represent the variability of relational structure observed in real networks. Inspired by concepts in boosting and bandit learning approaches, LBGA is designed to handle aggregations of noisy and disparate data sources, therefore marking a departure from methods that assume overlap and usefulness among all data sources considered.

LBGA also simplifies the task of designing a graph aggregation algorithm into designing a principled quality measure q and global event A . Doing so allows us to connect the utility of the graph representation to the application

of interest. As a byproduct, we conjecture that statistics produced by LBGA provide information about the utility of the graph source, such as the level of noise. Such information can be used to improve the data collection process and partially mitigate the effects noise before it propagates to subsequent analysis. We gave evidence for the resistance of our framework to noise by running it on datasets with various known levels of noise, and observed the resulting matching weight distributions.

Another primary concern for future work is to analyze the utility of our framework with respect to other graph applications, as well as to present a more thorough comparison of LBGA with existing multigraph clustering algorithms. We believe link prediction and label propagation to be ripe candidates, and there is an established similarity metric for the former of Adamic and Adar[2]. Our preliminary experiments have shown qualitatively different graph structure when using Adamic/Adar in place of *consistentNO*, but further research is necessary to fully understand the reason for it, and whether it corresponds to a qualitatively better analysis (in light of the wealth of literature on metrics for link prediction, we conjecture that it does).⁶

Additional directions include a more thorough stability analysis of LBGA, exploring the modifications we have suggested to improve scalability, and to prove further theoretical results as described earlier.

8 Acknowledgements

We thank Vineet Mehta for providing us with the Enron topic model data, and for his many helpful discussions.

References

- [1] Enron email dataset. <https://www.cs.cmu.edu/~enron/>.
- [2] L. Adamic and E. Adar. Friends and neighbors on the web. *Social Networks*, 25:211–230, 2001.
- [3] C.C. Aggarwal, Y. Xie, and P.S. Yu. Towards community detection in locally heterogeneous networks. In *SDM*, pages 391–402. SIAM, 2011.
- [4] S. Arora, E. Hazan, and S. Kale. The multiplicative weights update method: a meta-algorithm and applications. *Theory of Computing*, 8(1):121–164, 2012.
- [5] M. Balcan and A. Blum. *On a theory of kernels as similarity functions*. Manuscript, 2006.

⁶There is also a technical barrier to using Adamic/Adar directly, it is not a $[0,1]$ -valued function.

- [6] M. Balcan, A. Blum, and N. Srebro. A theory of learning with similarity functions. *Machine Learning*, 72(1-2):89–112, 2008.
- [7] M. Berlingerio, M. Coscia, and F. Giannotti. Finding redundant and complementary communities in multidimensional networks. In *CIKM*, pages 2181–2184, 2011.
- [8] S. Bubeck and N. Cesa-Bianchi. Regret analysis of stochastic and non-stochastic multi-armed bandit problems. *Foundations and Trends in Machine Learning*, 5(1):1–122, 2012.
- [9] S. Bubeck, R. Munos, and G. Stoltz. Pure exploration in multi-armed bandits problems. In *ALT*, pages 23–37, 2009.
- [10] R. Caceres, K. Carter, and J. Kun. A boosting approach to learning graph representations. In *Proc. SIAM Int. Conf. Data Mining, Workshop on Mining Networks and Graphs*, 2014. To appear.
- [11] R. S. Caceres, T. Y. Berger-Wolf, and R. Grossman. Temporal scale of processes in dynamic networks. In *ICDM Workshops*, pages 925–932, 2011.
- [12] D. Cai, Z. Shao, X. He, X. Yan, and Jiawei Han. Community mining from multi-relational networks. In *PKDD*, pages 445–452, 2005.
- [13] L. Danon, J. Duch, A. Diaz-Guilera, and A. Arenas. Comparing community structure identification. *J. Stat. Mech.*, 2005:P09008, 2005.
- [14] L. R. Dice. Measures of the amount of ecologic association between species. *Ecology*, 26(3):297–302, July 1945.
- [15] N. Eagle and A. Pentland. Reality mining: sensing complex social systems. *Personal and Ubiquitous Computing*, 10(4):255–268, 2006.
- [16] P. Erdős and A. Rényi. On random graphs, I. *Publicationes Mathematicae (Debrecen)*, 6:290–297, 1959.
- [17] S. Fortunato and M. Barthélémy. Resolution limit in community detection. *Proceedings of the National Academy of Sciences*, 104(1):36–41, 2007.
- [18] B. Gallagher, H. Tong, T. Eliassi-Rad, and C. Faloutsos. Using ghost edges for classification in sparsely labeled networks. In *Proc. (14th) ACM SIGKDD Inter. Conf. on Knowledge Discovery and Data Mining*, pages 256–264. ACM, 2008.
- [19] L. Getoor and C. P. Diehl. Link mining: a survey. *ACM SIGKDD Explorations Newsletter*, 7(2):3–12, 2005.
- [20] E. Gilbert and K. Karahalios. Predicting tie strength with social media. In *Proc. of SIGCHI Conf. on Human Factors in Computing Systems*, CHI ’09, pages 211–220, New York, NY, USA, 2009. ACM.

- [21] D. F. Gleich and C. Seshadhri. Vertex neighborhoods, low conductance cuts, and good seeds for local community methods. In *Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '12, pages 597–605, New York, NY, USA, 2012. ACM.
- [22] P. Jaccard. The distribution of the flora in the alpine zone. *New Phytologist*, 11(2):37–50, February 1912.
- [23] B. Klimt and Y. Yang. The enron corpus: A new dataset for email classification research. In *ECML*, pages 217–226, 2004.
- [24] J. Leskovec, K. J. Lang, A. Dasgupta, and M. W. Mahoney. Statistical properties of community structure in large social and information networks. In *Proc. (17th) Inter. Conf. on World Wide Web*, WWW '08, pages 695–704, New York, NY, USA, 2008. ACM.
- [25] M. Ley. The dblp computer science bibliography: Evolution, research issues, perspectives. In *SPIRE*, pages 1–10, 2002.
- [26] A. McCallum, A. Corrada-Emmanuel, and X. Wang. The author-recipient-topic model for topic and role discovery in social networks, with application to enron and academic email. In *Workshop on Link Analysis, Counterterrorism and Security*, pages 33–44, Newport Beach, CA, 2005.
- [27] A. K. McCallum. MALLET: A Machine Learning for Language Toolkit. <http://mallet.cs.umass.edu>, 2002.
- [28] B. A. Miller and N. Arcolano. Spectral subgraph detection with corrupt observations. In *Proc. IEEE Int. Conf. Acoust., Speech and Signal Process.*, 2014. To appear.
- [29] P. J. Mucha, T. Richardson, K. Macon, M. A. Porter, and J. P. Onnela. Community structure in time-dependent, multiscale, and multiplex networks. *Science*, 328(5980):876–878, 2010.
- [30] R. R. Nadakuditi and M. E. J. Newman. Graph spectra and the detectability of community structure in networks. *CoRR*, abs/1205.1813, 2012.
- [31] J. Neville and D. Jensen. Leveraging relational autocorrelation with latent group models. In *Proc. (4th) Inter. Workshop on Multi-relational Mining*, pages 49–55. ACM, 2005.
- [32] M. E. J. Newman. Modularity and community structure in networks. *Proceedings of the National Academy of Sciences of the United States of America*, 103(23):8577–8696, 2006.
- [33] E. E. Papalexakis, L. Akoglu, and D. Ience. Do more views of a graph help? community detection and clustering in multi-graphs. In *FUSION*, pages 899–905, 2013.

- [34] P. Pons and M. Latapy. Computing communities in large networks using random walks. *Journal of Graph Algorithms and Applications*, 10(2):191–218, 2006.
- [35] M. Rocklin and A. Pinar. On clustering on graphs with multiple edge types. *Internet Mathematics*, 9(1):82–112, 2013.
- [36] R. A. Rossi, L. McDowell, D. W. Aha, and J. Neville. Transforming graph data for statistical relational learning. *J. Artif. Intell. Res. (JAIR)*, 45:363–441, 2012.
- [37] R. E. Schapire. The strength of weak learnability. *Machine Learning*, 5:197–227, 1990.
- [38] L. Tang, X. Wang, and H. Liu. Community detection via heterogeneous interaction analysis. *Data Min. Knowl. Discov.*, 25(1):1–33, 2012.
- [39] W. Tang, Z. Lu, and I. S. Dhillon. Clustering with multiple graphs. In *Proc. (2009) IEEE Int. Conf. on Data Mining, ICDM '09*, pages 1016–1021, Washington, DC, USA, 2009. IEEE Computer Society.
- [40] Y. Wang and G. Wong. Stochastic Block Models for Directed Graphs. *Journal of the American Statistical Association*, 82(397):8–19, 1987.
- [41] R. Xiang, J. Neville, and M. Rogati. Modeling relationship strength in online social networks. In *Proc. (19th) Int. Conf. on World Wide Web, WWW '10*, pages 981–990, New York, NY, USA, 2010. ACM.